

Analysis of Algorithms

V. Adamchik

CSCI 570

Lecture 13

University of Southern California

Spring 2023

Approximation Algorithms

The Design of Approximation Algorithms,

D. Williamson and D. Shmoys, Cambridge University Press, 2011.

Exam - II

Date: Thursday April 27

Time: starts at 5pm

Length: 2 hrs and 20 mins

Locations: TBA

Practice: posted

TA Review: April 25 and April 26

Closed book and notes

Follow the Honor Code (obey all rules for taking exams)

No internet searching

No talking to each other (chat, phone, messenger)

Exam - II

True/False

Multiple Choice

Written Problems:

- Network Flow (max flow, circulation)
- Linear Programming (standard and dual forms)
- NP Completeness (a proof by reduction)
- Approximations

Approximation Algorithms



"Although this may seem a paradox, all exact science is dominated by the idea of approximation",
Bertrand Russell (1872-1970)

Suppose we are given an NP-hard problem to solve.

Can we develop *a polynomial-time algorithm* that produces a "good enough" solution?


An algorithm that returns near-optimal solutions is called an *approximation algorithm*.

Graph Coloring

Given a graph $G=(V,E)$, find the **minimum** number of colors required to color vertices, so no two adjacent vertices have the same color.

This is NP-hard problem.

Let us develop a solution that is close enough to the optimal solution.



Greedy Approximation

Given $G = (V, E)$ with n vertices.

Use the integers $\{1, 2, 3, \dots, n\}$ to represent colors.

Order vertices by degree in descending order.


Color the first vertex (highest degree) with color 1.

Go down the vertex list and color every vertex not adjacent to it with color 1.

Remove all colored vertices from the list.

Repeat for uncolored vertices with color 2.

Example



Greedy Order: H, K, D, G, I, J, A, B, E, F, C


Formal Definition

Let P be a minimization (convex) problem, and I be an instance of P .

Let $ALG(I)$ be a solution returned by an algorithm.

Let $OPT(I)$ be the optimal solution.

Then $ALG(I)$ is said to be a α -approximation algorithm for some $\alpha > 1$, if for $ALG(I) \leq \alpha \cdot OPT(I)$.



These notions allow us to circumvent NP-hardness by designing polynomial-time algorithm with worst-case guarantees!


Maximization Problem

Let P be a maximization (convex) problem, and I be an instance of P .

Let $ALG(I)$ be a solution returned by an algorithm.

Let $OPT(I)$ be the optimal solution.

Then $ALG(I)$ is said to be a α -*approximation* algorithm for some $0 < \alpha < 1$, if for $ALG(I) \geq \alpha \cdot OPT(I)$.



Vertex cover

Given $G=(V,E)$, find **the smallest** $S \subset V$ s.t. every edge is incident on a vertex in S .

Let us design a randomized approximation algorithm for vertex cover.

APPROX-VC($G=(V,E)$)

$VC \leftarrow \emptyset$ (vertex cover)

$E' \leftarrow E(G)$



while $E' \neq \emptyset$ do

 let (u, v) be an arbitrary edge of E'


$VC \leftarrow VC \cup \{u, v\}$

 remove every edge in E' incident on u or v

return VC



Example




Vertex Cover

Lemma. Let M be a matching in G , and VC be a vertex cover, then

$$|M| \leq |VC|$$

Proof.



2-Approximation Vertex Cover

Approx-VC(G):

M - maximal matching on G


VC - take both endpoints of edges in M

Return VC

Theorem. Let $\text{OPT}(G)$ be the size of the optimal vertex cover and $\text{VC} = \text{ApproxVC}(G)$. Then $|\text{VC}| \leq 2 \cdot \text{OPT}(G)$.

Proof.

Can we do better than 2-Approximation?



Traveling Salesman Problem

The traveling salesman problem consists of a salesman and a set of cities (a complete weighted graph). The salesman has to visit each one of the cities starting from a certain one (e.g. the hometown) and returning to the same city. The challenge of the problem is that the traveling salesman wants to minimize the total length of the trip.

Theorem.

The traveling salesman problem is NP-hard.

Proof

We prove that TSP is NP-hard by reduction from Hamiltonian Cycle (HC): $HC \leq_p TSP$

Solving TSP

TSP: given an undirected complete graph $G=(V,E)$ with edge cost $c:E\rightarrow\mathbb{R}^+$, find a min cost Hamiltonian cycle.

We will consider a **metric** TSP.

In the metric TSP, edge costs have to satisfy the triangle inequality, i.e. for any three vertices x, y, z ,




$$c(x,z) \leq c(x,y) + c(y,z).$$

We develop approximation greedy algorithm.

Approximation Greedy Algorithm

Approx-TSP(G):

- 1) Find a MST of G (complete graph)
- 2) Create a cycle by doubling edges
- 3) Remove double visited edges



Approximation TSP

Theorem. Approx-TSP is a 2-approximation algorithm for a metric TSP.

Proof.

$$|\text{Approx-TSP}| \leq 2 \cdot |\text{MST}| \leq 2 \cdot |\text{OPT}|$$

↑
triangle inequality

↑
we can get a spanning
tree from HC by
removing edges

Christofides Algorithm

Observe that a factor 2 in the approximation ratio is due to doubling edges.

But any graph with even degrees vertices has a Eulerian cycle.
A Eulerian path visits each edge exactly once

Intuition: we have to add edges only between odd degree vertices

Christofides Algorithm

Theorem (1976)(without proof).

Christofides is $3/2$ approximation for Metric TSP

The algorithm has been known for over 50 years and yet no improvements have been made since its discovery.

General TSP

Theorem: *If $P \neq NP$, then for $\forall \alpha > 1$ there is NO a polynomial time α -approximation of *general TSP*.*

Proof.

Suppose for contradiction that such an α -approximation algorithm exists. We will use this algorithm to solve the Hamiltonian Cycle problem.

Start with $G=(V,E)$ and create a new complete graph G' with the cost function

$$c(u,v) = 1, \text{ if } (u,v) \in E$$

$$c(u,v) = \alpha \cdot V + 1, \text{ otherwise}$$

General TSP

Proof (continue)

$$c(u,v) = 1, \text{ if } (u,v) \in E$$

$$c(u,v) = \alpha \cdot V + 1, \text{ otherwise}$$

If G has HC , then $|\text{TSP}(G')| = V$.

If G has no HC , then

$$|\text{TSP}(G')| \geq (V - 1) + \alpha \cdot V + 1 = V + \alpha \cdot V \geq \alpha \cdot V$$

Since the $|\text{TSP}|$ differs by a factor α , our approximation algorithm can be able to distinguish between two cases, thus decides if G has a ham-cycle in polynomial time. Contradiction.

Bin Packing

You have an infinite supply of bins, each of which can hold M maximum weight. You also have n objects, each of which has a weight w_i (w_i is at most M). Our goal is to partition the objects into bins, such that we use as few bins as possible.

This problem in general is NP-hard.

Devise a 2-approximation algorithm.

First-fit approach


1. process items in the original order
2. if an item doesn't fit the first bin, put it into the next bin
3. if an item does not fit into any bins, open a new bin

Runtime-?

Example,

items = {40, 20, 35, 15, 25, 5, 30, 10}

$M = 45$



Proof: a 2-approximation

Suppose we use m bins.

Let m^* denote the optimal number of bins.

Clearly, $m^* \geq (\sum w_i)/M$,

On the other hand, $(\sum w_i)/M > (m-1)/2$

This follows from the fact that at least $(m-1)$ bins are more than half full.


Because, if there are two bins less than half full, items in the second bin will be placed into the first by our algorithm.

We have showed, $m^* \geq (\sum w_i)/M > (m-1)/2$

Comparing the left and right hand sides, we derive

$$2m^* > m-1 \quad \text{or,} \quad m \leq 2m^*$$

Developing a lower bound ($M=1$)




Insert $2N$ items: $\frac{1}{2} - \varepsilon, \dots, \frac{1}{2} - \varepsilon, \frac{1}{2} + \varepsilon, \dots, \frac{1}{2} + \varepsilon$

OPT: N bins

FF: $N/2$ bins + N bins = $1.5 N$

$FF \geq 1.5 OPT$

Developing a lower bound ($M = 1$)




$3N$ items: $1/7 + \epsilon, \dots, 1/7 + \epsilon, 1/3 + \epsilon, \dots, 1/3 + \epsilon, 1/2 + \epsilon, \dots, 1/2 + \epsilon$

OPT: N bins

FF: $N/6 + N/2 + N$ bins = $5/3 N$

$FF \geq 1.66 \text{ OPT}$

Developing a lower bound



Observe

$$1 - (1/2 + 1/3 + 1/7) = 1/42$$

4N items: $1/43 + \varepsilon, \dots, 1/43 + \varepsilon,$
 $1/7 + \varepsilon, \dots, 1/7 + \varepsilon,$
 $1/3 + \varepsilon, \dots, 1/3 + \varepsilon,$
 $\frac{1}{2} + \varepsilon, \dots, \frac{1}{2} + \varepsilon$


OPT: N bins

FF: $N/42 + N/6 + N/2 + N = 71/42 N$

$FF \geq 1.6905 \text{ OPT}$

Developing a lower bound

Observe, $1 - (1/2 + 1/3 + 1/7 + 1/43) = 1/1806$



$1/1807 + \epsilon$
 $1/43 + \epsilon$

FF approximation

=

$N(1 + 1/2 + 1/6 + 1/42 + 1/1806 + \dots)$

= (1.691..) OPT

Theorem (1976)
*This is a 17/10-
approximation
algorithm.*

In the limit, the ratio is 1.7

OPT: N bins

The formal proof is quite complex.

Sorted First-fit approach


1. sort the input in descending order
2. apply the first-fit algorithm

Theorem.
*This is a $11/9$ -
approximation
algorithm.*

Example,

items = {40, 20, 35, 15, 25, 5, 30, 10}

$M = 45$



Discussion Problem 1

Suppose you are given a set A of positive integers a_1, a_2, \dots, a_n and a positive integer B . A subset $S \subseteq A$ is called *feasible* if the sum of the numbers in S does not exceed B . You would like to select a feasible subset $S \subseteq A$ whose total sum is as large as possible. Give a linear-time algorithm that finds a feasible set $S \subseteq A$ whose total sum is at least half as large as the maximum total sum of any feasible set. You may assume that $a_i \leq B$.

Example: If $A = \{8, 2, 4, 3, 5\}$ and $B = 11$ then the optimal solution is the subset $S = \{8, 3\}$.

Load Balancing

There are m identical machines and n jobs, where each job j has processing time t_j . Job j must run contiguously on one machine. A machine can process at most one job at a time.

Let $J(i)$ be the subset of jobs assigned to machine i .

The **load** of machine i is $L_i = \sum_{j \in J(i)} t_j$.

The **makespan** is the maximum load on any machine $L = \max L_i$.

Load balancing problem:

assign each job to a machine to minimize makespan.